

Analysis of Positional Encoding and Attention Mechanisms in Disaster-Related Tweet Classification




vs




Javier Rodriguez & Edgar Leon

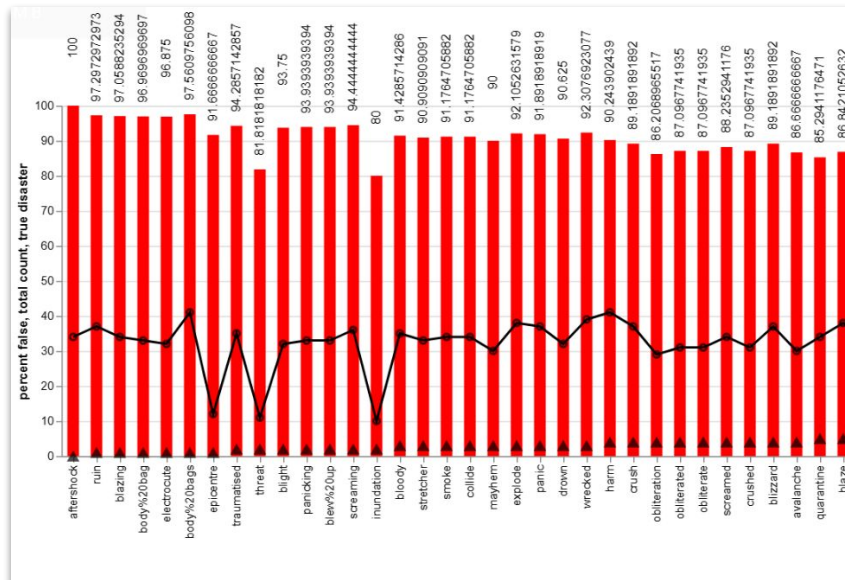

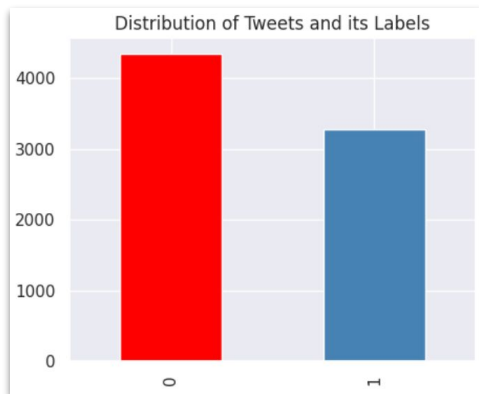
Disaster Tweets Input Data



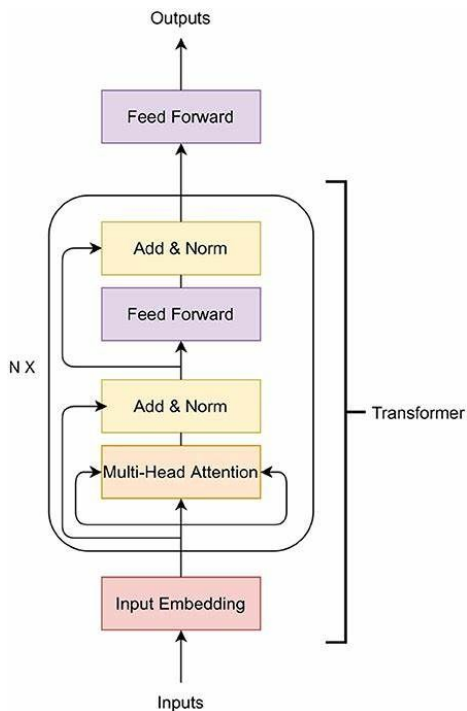
Disaster Tweets Input Data



| | id | keyword | location | text | target |
|-----|-----|------------|-------------------------|---------------------------------------------------|--------|
| 100 | 144 | accident | UK | .@NorwayMFA #Bahrain police had previously die... | 1 |
| 101 | 145 | accident | Nairobi, Kenya | I still have not heard Church Leaders of Kenya... | 0 |
| 102 | 146 | aftershock | Instagram - @heyimginog | @afterShock_DeLo scuf ps live and the game... cya | 0 |
| 103 | 149 | aftershock | 304 | 'The man who can drive himself further once th... | 0 |
| 104 | 151 | aftershock | Switzerland | 320 [IR] ICEMOON [AFTERSHOCK] http://t.co/yN... | 0 |



NMT Classifier



Define the components

The embedding and positional encoding layer

Add and normalize

The base attention layer

The global self attention layer

The feed forward network

The encoder layer

The encoder

The Transformer Classifier

The Classifier Transformer Architecture Along with its Components

This classifier model is based on last hidden state

```
def positional_encoding(length, depth):
    depth = depth/2

    positions = np.arange(length)[::, np.newaxis] # (seq, 1)
    depths = np.arange(depth)[np.newaxis, :]/depth # (1, depth)

    angle_rates = 1 / (10000**depths) # (1, depth)
    angle_rads = positions * angle_rates # (pos, depth)

    pos_encoding = np.concatenate(
        [np.sin(angle_rads), np.cos(angle_rads)],
        axis=-1)

    return tf.cast(pos_encoding, dtype=tf.float32)

"""
The position encoding function is a stack of sines and cosines that vibrate at different
along the depth of the embedding vector. They vibrate across the position axis.
"""

class PositionalEmbedding(tf.keras.layers.Layer):
    def __init__(self, vocab_size, d_model):
        super().__init__()
        self.d_model = d_model
        self.embedding = tf.keras.layers.Embedding(vocab_size, d_model, mask_zero=True)
        self.pos_encoding = positional_encoding(length=2048, depth=d_model)
        #the line below is new
        #self.pos_encoding = tf.cast(self.pos_encoding, tf.float64)

    def compute_mask(self, *args, **kwargs):
        return self.embedding.compute_mask(*args, **kwargs)

    def call(self, x):
        #x, _ = x #ignore attn weights
        length = tf.shape(x)[1]
        x = self.embedding(x)
        # This factor sets the relative scale of the embedding and positional_encoding.
        x *= tf.math.sqrt(tf.cast(self.d_model, tf.float32))
        #x = tf.math.sqrt(tf.cast(self.d_model, tf.float64))
        x = x + self.pos_encoding[tf.newaxis, :length, :]
        #print("positional embedding: shape and type:")
        #the casting is new
        #x = tf.cast(x, tf.float64)
        #print(x.dtype)
        #print(tf.shape(x)[1])
        return x


"""The base attention layer"""
class BaseAttention(tf.keras.layers.Layer):
    def __init__(self, **kwargs):
        super().__init__()
        self.mha = tf.keras.layers.MultiHeadAttention(**kwargs)
        self.layernorm = tf.keras.layers.LayerNormalization()
```

First results on raw data




| | id | keyword | location | text | target |
|-----|-----|------------|-------------------------|---------------------------------------------------|--------|
| 100 | 144 | accident | UK | .@NorwayMFA #Bahrain police had previously die... | 1 |
| 101 | 145 | accident | Nairobi, Kenya | I still have not heard Church Leaders of Kenya... | 0 |
| 102 | 146 | aftershock | Instagram - @heyimginog | @afterShock_DeLo scuf ps live and the game... cya | 0 |
| 103 | 149 | aftershock | 304 | 'The man who can drive himself further once th... | 0 |
| 104 | 151 | aftershock | Switzerland | 320 [IR] ICEMOON [AFTERSHOCK] http://t.co/yN... | 0 |

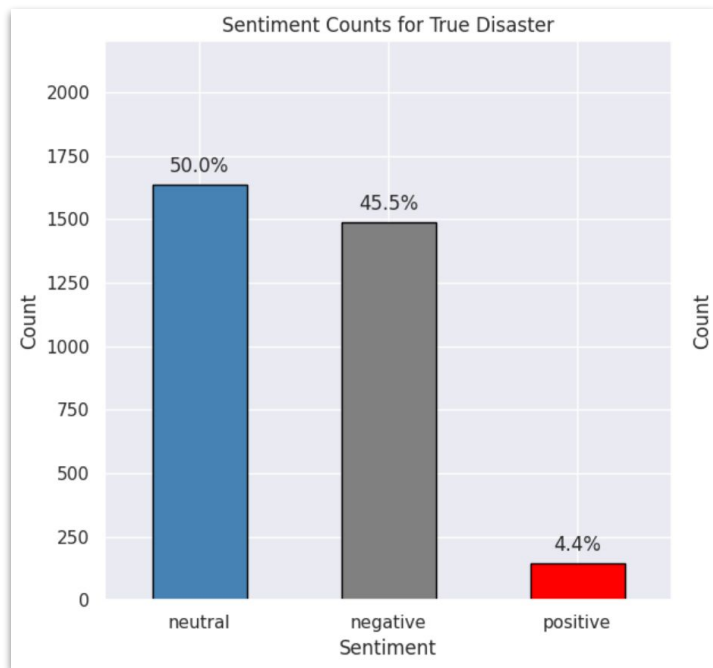
Epoch 10/10
84/84 [=====] - 71s 851ms/step - loss: 0.3064 - accuracy: 0.8760 - val_loss: 0.5612 - val_accuracy: 0.7806



val accuracy: 0.78




Input Data EDA – content analysis & enhancement




Tweets enhanced by:

- Location
- Keyword
- Url
- Sentiment
- topic
- Irony
- emotion



location ablaze **bbcmt**d wholesale markets ablaze <http://t.co/lhyxeohy6c> neutral news_&_social_concern irony ang



Data Wrangling Final Results

```
[ ] unknown_words, known_words = find_unknown_words(sentences, word_index)
```

```
# Call the function and print the unknown words
print(f"Words found by the tokenizer: {len(known_words)}")
print(f"Words not found by the tokenizer: {len(unknown_words)}")
```

```
Words found by the tokenizer: 17980
Words not found by the tokenizer: 9
```

```
[ ] for i in range(len(unknown_words)):
    print(list(unknown_words)[i])
```

```
`bbcnews
^oo^
^ag
^mp
\\
3\30a
destruction\s
didn't
^sj
```

Creation of Embeddings

-The "fasttext_english_twitter_100d.vec" file is a set of word vectors trained using the FastText library.

FastText is a library developed by Facebook's AI Research (FAIR) lab. It's dedicated to text classification and learning word representations. FastText models can be trained on more than a billion words on any multicore CPU in less than a few minutes.


Here's a general process of how FastText trains word vectors:

Corpus Selection: FastText trains word vectors on large text corpora. For example, FastText provides pre-trained word vectors trained, in this case, a collection of English tweets.


```
[ ] embedding_df.shape
```

```
(17984, 100)
```



Clean Data and Creation of Embeddings



Clean and enhanced data
no embeddings: 0.79



Raw data no embeddings:
0.78

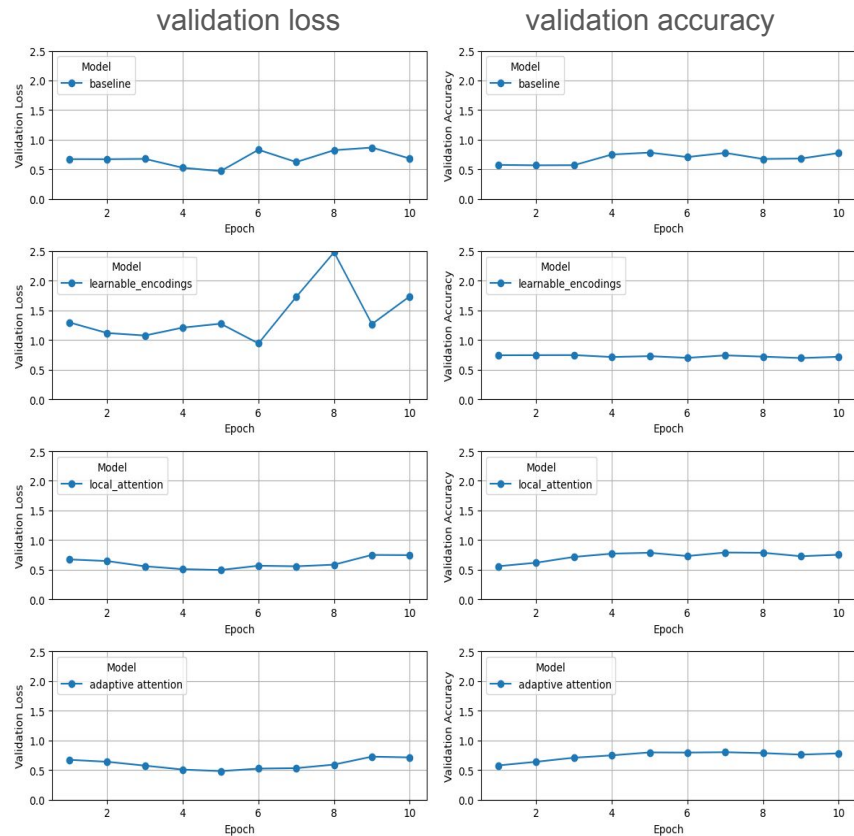


Clean and enhanced data
with embeddings: 0.84


Next explore other architecture options

Model Comparison


| | |
|-------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Model 1: Baseline | <p>Validation Accuracy: Peaked at 78.89% but fluctuated towards the end, settling at 77.76%.</p> |
| Model 2: learnable encodings | <p>Training Accuracy: Extremely high from the start, reaching 99.42%.</p> <p>Validation Loss: Started high and increased over epochs, indicating potential overfitting</p> |
| Model 3: local attention | <p>Training Accuracy: Improved consistently from 48.71% to 96.58%.</p> <p>Validation Accuracy: Increased to a peak of 80.04%, then fluctuated.</p> <p>Validation Loss: Started low, decreased to a point, then increased significantly in later epochs.</p> |
| Model 4: adaptive attention | <p>Training Accuracy: It shows robust improvement in training accuracy, the highest peak in validation accuracy, and while it does show an increase in validation loss indicating some overfitting.</p> |



Positional Encoding with Twitter Pre-Trained Embeddings



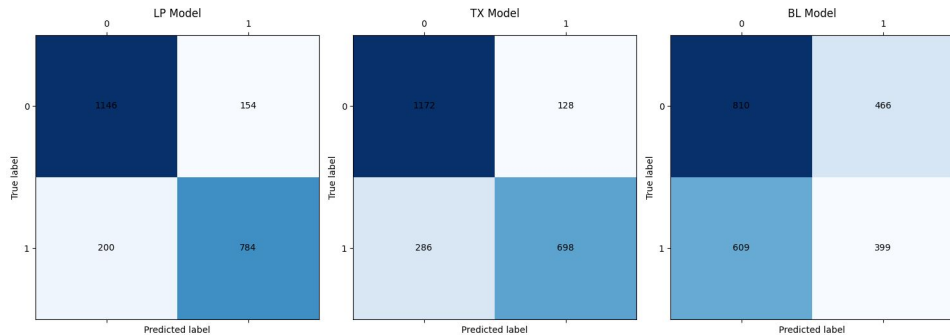
Baseline Transformer: 0.84



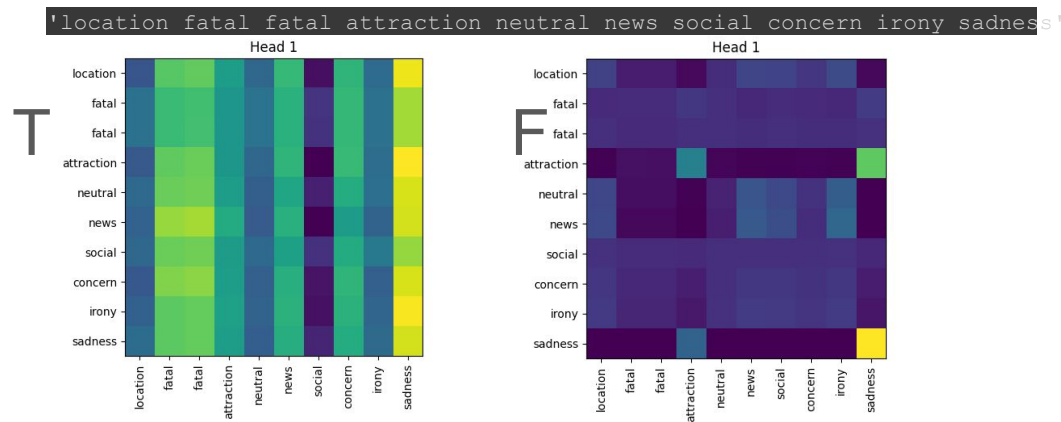
Learnable Positional Encoding: 0.82

LPE vs Baseline Transformers Analysis

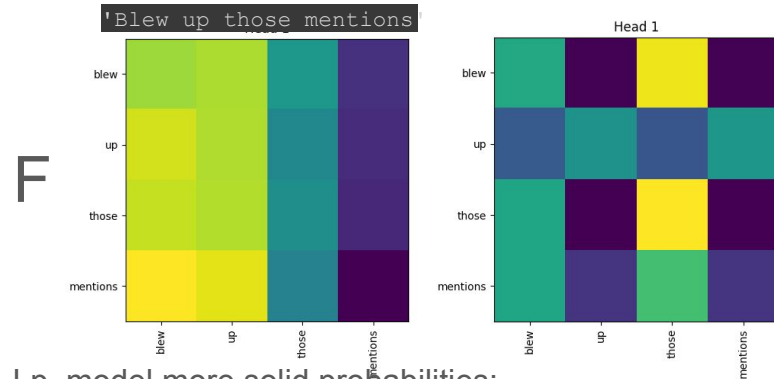
Lp_model more balanced:



Lp_model pays better attention:



Lp_model more nuanced understanding:



Lp_model more solid probabilities:




Analysis Updates to Embeddings

Lp_model fine tuned each word more:



tx_model



lp_model

On average lp_model changed embeddings more (0.43 vs 0.22):

```
tx_disaster_words = top_changed_embeddings(loaded_embeddings, tx_embeddings_after_training, word_index)
tx_disaster_words_list = [item[0] for item in tx_disaster_words]
tx_disaster_words

[('suicide', 0.2542814085893426),
 ('hiroshima', 0.236663998735998),
 ('derailment', 0.23234185947136963),
 ('positive', 0.2134087088705059),
 ('california', 0.21303308061344782),
 ('bombing', 0.21064468542664885),
 ('debris', 0.2106077369589732),
 ('wildfire', 0.20714493198224423),
 ('killed', 0.200944669891897),
 ('northern', 0.19962536227301594)]

lp_disaster_words = top_changed_embeddings(loaded_embeddings, lp_embeddings_after_training, word_index)
lp_disaster_words_list = [item[0] for item in lp_disaster_words]
lp_disaster_words

[('hiroshima', 0.496638707030925),
 ('spill', 0.4849229581973467),
 ('suicide', 0.448507128728351),
 ('bomb', 0.4369772545959775),
 ('derailment', 0.4233027778885434),
 ('northern', 0.418132935104703),
 ('california', 0.4133995936162507),
 ('atomic', 0.40671896164405086),
 ('killed', 0.39854748212169594),
 ('wreckage', 0.3949224709598692)]
```

Challenge #3 Type of Disaster? Mislabeled as True



Lp_model did 82% while tx_model did 84%, we attribute it to mislabeled data:

| index | id | keyword | location | text | target | comma |
|-------|------|---------|-------------|--------|---------------------------------------------------|-------|
| 108 | 229 | 328 | annihilated | NaN | Ready to get annihilated for the BUCS game | 1 |
| 281 | 753 | 1085 | blew%20up | NaN | @BenKin97 @Mili_5499 remember when u were up l... | 1 |
| 709 | 1831 | 2632 | crashed | London | This guy bought my car on Tuesday police knock... | 1 |
| 872 | 2310 | 3318 | demolished | NaN | Got my first gamer troll I just demolished a k... | 1 |
| 1154 | 2946 | 4237 | drowned | NaN | I got drowned like 5 times in the damn game to... | 1 |

Given better data, lp_model should do better:

| | | | | sentence | length |
|------|----------|--------------------------|------------|--------------------|--------|
| 28 | location | coool | positive | news social con... | 61 |
| 24 | location | looooool | neutral | news social con... | 65 |
| 21 | location | london is cool | positive | news so... | 69 |
| 17 | location | summer is lovely | positive | news s... | 70 |
| 6705 | location | thunder thunder | neutral | news socia... | 71 |
| 18 | location | my car is so fast | positive | news ... | 71 |
| 3749 | location | fire i see fire | neutral | news socia... | 71 |
| 1882 | location | crushed crushed | neutral | news socia... | 74 |
| 3608 | location | fatal fatal | attraction | neutral news ... | 77 |
| 26 | location | was in nyc last week | neutral | news ... | 77 |
| 2496 | location | desolate eggs | desolate | negative news... | 78 |
| 4922 | location | mayhem mayhem is | beautiful | positive ... | 78 |
| 3688 | location | fatality fatality | https | neutral new... | 79 |
| 4184 | location | hazard get that hazard | pay | neutral n... | 79 |
| 4092 | location | hail hail pic %00 | https | positive ne... | 80 |
| 3751 | location | fire im on fire | weblink | neutral ... | 80 |
| 7448 | location | wounds 11 puncture | wounds | neutral ne... | 80 |
| 6224 | location | smoke smoke eat sleep | neutral | news ... | 80 |
| 388 | location | arson the sound of arson | negative | ne... | 81 |
| 1886 | location | crushed crushed it | https | positive n... | 83 |
| 823 | location | blizzard blizzard | gamin ight | neutral ... | 83 |
| 3717 | location | fear my worst fear | https | negative ... | 83 |

Future Work & Demo

EmergencyGPT


Disclaimer: This application is not a substitute for 911.

Which emergency would you like to report?

fire

AI-sistance


Our attention to you

 Please report to us TRUE emergencies only.

Please enter your question

Appendix

Challenge #4 Wide Range of Twitter Users



Age Difference



Cultural Difference



Model Comparison





training loss

training accuracy





validation loss

validation accuracy





baseline



learnable encodings



local attention



adaptive attention

